REAL-TIME RAILWAY TRACK MONITORING AND DEFECT DETECTION

This is the part of the system that makes the real-time tracking and evaluating of railway tracks with cameras on inspection vehicles automatic. The cameras constantly take pictures and video frames during the train's track journey. The system uses image processing and computer vision techniques for the captured frames to detect rail cracks, broken rails, damaged sleepers, missing or broken fasteners, surface defects, track deformation and abnormal objects or obstacles. This cuts the need for inspection and allows for ongoing monitoring of the railway system's condition.

The captured frames are then preprocessed using image resizing, noise filtering, gray level conversion, contrast enhancement, sharpening, and edge detection and illumination correction to increase the quality of detection. The system can be used in railway environments such as vehicle vibration, motion blur, weak light, shadows, rust and complex backgrounds. Object detection and classification methods are used to identify and classify the detected objects, and the use of classification based on the confidence level can distinguish the real defects from normal railway parts and environmental factors.

The title can meet the performance requirements for detection accuracy of ≥92%, a processing rate of ≥15 frames per second (fps), a processing latency of ≤100 ms per frame and a false-positive rate of ≤5%. The detection pipeline is optimized for computational resources on an edge device allowing images to be processed locally without depending completely on cloud computing. If the defect is picked up, it is entered in the system with data on the type of defect, its location, its level of confidence, when it was detected, and its severity. Critical defects can generate alerts, to railway monitoring personnel while recorded inspection data can be used for further analysis and maintenance planning.

This title will deliver a rapid, reliable railway track inspection mechanism which will aid early detection of harmful track conditions. It assists railway operators in conducting inspection tasks and carrying out maintenance, reduce manpower effort in inspection, reduce false alarm and enhance the reliability of railway infrastructure. The module integrates OpenCV-based image processing, computer vision and real-time video analysis and edge computing to facilitate monitoring and plays a significant role in railway safety monitoring and accident prevention.

Tools used:

- Python
- OpenCV
- YOLO
- NumPy
- PyTorch
- Roboflow
- NVIDIA Jetson / Raspberry Pi

Pseudo Code:

```
config.py — Global Configuration for Railway Track Monitoring System
Central configuration hub for all system parameters including model paths,
detection thresholds, performance targets, class definitions, and visual settings.
"""
import os
from pathlib import Path
BASE_DIR = Path(__file__).resolve().parent
DATA_DIR = BASE_DIR / "data" / "dataset"
MODEL_DIR = BASE_DIR / "model" / "weights"
OUTPUT_DIR = BASE_DIR / "outputs"
FIGURES_DIR = OUTPUT_DIR / "figures"
REPORTS_DIR = OUTPUT_DIR / "reports"
for _dir in [DATA_DIR, MODEL_DIR, OUTPUT_DIR, FIGURES_DIR, REPORTS_DIR]:
    _dir.mkdir(parents=True, exist_ok=True)
 Dataset Configuration
DATASET_YAML = DATA_DIR / "dataset.yaml"
IMG_DIR = DATA_DIR / "images"
LABEL_DIR = DATA_DIR / "labels"
NUM_SYNTHETIC_IMAGES = 600          # Total synthetic images to generate
TRAIN_SPLIT = 0.8                   # 80% train, 20% val
IMG_SIZE = 640                      # YOLO input resolution (px)
CLASS_NAMES = [
    "crack",           # 0 — Surface / rail-head cracks
    "broken_fastener", # 1 — Missing or fractured rail clips/bolts
    "surface_defect",  # 2 — Corrosion, spalling, gauge damage
    "foreign_object",  # 3 — Debris, stones, vegetation on track
]
NUM_CLASSES = len(CLASS_NAMES)
CLASS_SEVERITY = {
```

```
    "crack":          "CRITICAL",
    "broken_fastener": "HIGH",
    "surface_defect":  "MEDIUM",
    "foreign_object":  "HIGH",
}
CLASS_COLORS = {
    "crack":          (0,   30, 220),   # Red
    "broken_fastener": (0,  140, 255),   # Orange
    "surface_defect":  (0,  200, 200),   # Yellow
    "foreign_object":  (220, 80,   0),   # Blue
}
YOLO_BASE_MODEL = "yolov8n.pt"          # Nano — fastest for edge deployment
TRAINED_MODEL_PT  = MODEL_DIR / "best.pt"
TRAINED_MODEL_ONNX = MODEL_DIR / "best.onnx"
TRAIN_EPOCHS   = 30
TRAIN_BATCH    = 16
TRAIN_WORKERS  = 4
TRAIN_PATIENCE = 10                  # Early stopping patience
CONFIDENCE_THRESHOLD = 0.45          # Minimum detection confidence
IOU_THRESHOLD        = 0.45          # NMS IoU threshold
MAX_DETECTIONS       = 50            # Max detections per frame
# Temporal false-positive suppression
TEMPORAL_WINDOW      = 3             # Consecutive frames to check
TEMPORAL_MIN_HITS    = 2             # Must appear in ≥ this many frames
# ROI (Region of Interest) — fraction of frame height to consider track region
ROI_Y_START = 0.20                   # Top boundary (20% from top)
ROI_Y_END   = 0.90                   # Bottom boundary (90% from top)
TARGET_ACCURACY_PCT  = 92.0          # mAP@0.5 ≥ 92%
TARGET_LATENCY_MS    = 100.0         # Per-frame latency ≤ 100 ms
TARGET_FPS           = 15.0          # Throughput ≥ 15 FPS
```

```
TARGET_FP_RATE_PCT   = 5.0          # False-positive rate ≤ 5%
CLAHE_CLIP_LIMIT    = 2.0           # CLAHE contrast limit
CLAHE_TILE_SIZE     = (8, 8)        # CLAHE tile grid size
DENOISE_STRENGTH    = 7             # fastNlMeans denoising strength (h param)
DENOISE_TEMPLATE    = 7             # Template patch size
DENOISE_SEARCH      = 21            # Search window size
SHADOW_SAT_THRESH   = 40            # Shadow removal saturation threshold
OVERLAY_ALPHA       = 0.35          # Bounding box fill transparency
FONT_SCALE          = 0.55
FONT_THICKNESS      = 2
BOX_THICKNESS       = 2
HUD_BG_ALPHA        = 0.6           # HUD background transparency
HUD_TEXT_COLOR      = (220, 220, 220) # Light grey for HUD text
# Alert flash colors (BGR)
ALERT_COLOR_CRITICAL = (0,   10, 200)
ALERT_COLOR_HIGH     = (0,  130, 255)
ALERT_COLOR_MEDIUM   = (0,  200, 200)
ALERT_COLOR_LOW      = (80, 200,  80)
ALERT_FLASH_FRAMES   = 8            # Frames to show alert flash
COMPARISON_METHODS = {
    "Edge + Hough Transform": {
        "accuracy": 88.0,
        "latency_ms": 120.0,
        "fps": 25.0,
        "fp_rate": 12.0,
        "edge_suitable": True,
        "decision": "REJECT — accuracy < 92%",
    },
    "SIFT Feature Matching": {
        "accuracy": 90.0,
```

```
    "latency_ms": 95.0,
    "fps": 18.0,
    "fp_rate": 8.5,
    "edge_suitable": False,
    "decision": "REJECT — accuracy < 92%",
},
"Viola–Jones": {
    "accuracy": 84.0,
    "latency_ms": 55.0,
    "fps": 30.0,
    "fp_rate": 15.0,
    "edge_suitable": True,
    "decision": "REJECT — accuracy < 92%",
},
"YOLOv8n (Ours)": {
    "accuracy": None,              # Filled by benchmarker at runtime
    "latency_ms": None,
    "fps": None,
    "fp_rate": None,
    "edge_suitable": True,
    "decision": "SELECTED",
},
"Deep Learning (EfficientDet)": {
    "accuracy": 97.0,
    "latency_ms": 130.0,
    "fps": 10.0,
    "fp_rate": 2.5,
    "edge_suitable": False,
    "decision": "REJECT — latency > 100 ms, not edge-suitable",
},
```

}

Main Code:

main.py — Railway Track Monitoring System: Entry Point

Real-Time Railway Track Defect Detection using YOLOv8 + OpenCV

Assignment: Computer Vision — Real-Time Railway Track Monitoring

Objective: Detect cracks, broken fasteners, surface defects, and
foreign objects with ≥92% accuracy, ≤100ms latency,
≥15 FPS, and ≤5% false positive rate.

Architecture Decision:

YOLO was selected over classical methods (Edge+Hough, SIFT, Viola–Jones)
and heavy deep learning models (EfficientDet) because it uniquely satisfies
ALL four performance constraints simultaneously while remaining deployable
on edge devices via ONNX/TensorRT optimization.

Quick Start:

```
python main.py                # Full benchmark + report (no GUI)
python main.py --demo         # Live OpenCV detection window
python main.py --train        # Generate dataset + train YOLOv8n
python main.py --report-only  # Generate report figures from existing results
```

For advanced options:

```
import sys
import time
import argparse
from pathlib import Path
def print_banner() -> None:
    banner = (
        "\n" +
        "=" * 66 + "\n" +
        "  REAL-TIME RAILWAY TRACK MONITORING SYSTEM\n" +
        "  YOLOv8 + OpenCV | Edge-Optimized Pipeline\n" +
        "=" * 66 + "\n" +
```

```
        "  Defect Classes : Crack | Broken Fastener | Surface | Foreign\n" +
        "  Requirements   : >=92% Acc | <=100ms Lat | >=15 FPS | <=5% FP\n" +
        "  Backend        : PyTorch -> ONNX -> TensorRT (edge deployment)\n" +
        "=" * 66 + "\n"
    )    print(banner)
def print_system_info() -> None:
    """Print environment and dependency information."""
    import platform
    print("── System Information
    print(f"  Python   : {sys.version.split()[0]}")
    print(f"  Platform : {platform.system()} {platform.release()}")
    try:
        import torch
        print(f"  PyTorch  : {torch.__version__}")
        cuda = torch.cuda.is_available()
        device = f"CUDA ({torch.cuda.get_device_name(0)})" if cuda else "CPU only"
        print(f"  CUDA     : {'Available' if cuda else 'Not available'}  [{device}]")
    except ImportError:
        print("  PyTorch  : NOT INSTALLED (run: pip install torch)")
    try:
        import cv2
        print(f"  OpenCV   : {cv2.__version__}")
    except ImportError:
        print("  OpenCV   : NOT INSTALLED (run: pip install opencv-python)")
    try:
        # pyrefly: ignore [missing-import]
        import ultralytics
        print(f"  Ultralytics : {ultralytics.__version__}")
    except ImportError:
        print("  Ultralytics : NOT INSTALLED (run: pip install ultralytics)")
```

```
    try:
        # pyrefly: ignore [missing-import]
        import onnxruntime as ort
        print(f"  OnnxRuntime : {ort.__version__}")
    except ImportError:
        print("  OnnxRuntime : NOT INSTALLED (run: pip install onnxruntime)")
def run_full_pipeline(
    num_benchmark_frames: int = 200,
    use_trained: bool = False,
    skip_report: bool = False,
    verbose: bool = True
) -> dict:
    """
    Execute the complete evaluation pipeline:
      1. Initialize detector (auto: ONNX → PyTorch → Mock)
      2. Run performance benchmark (200 synthetic frames with ground truth)
      3. Build method comparison table
      4. Generate 5 academic report figures
      5. Print summary table
    Returns:
        dict with benchmark metrics and report file paths.
    """
    from model.detector           import Detector
    from pipeline.preprocessor    import Preprocessor
    from pipeline.postprocessor   import PostProcessor
    from validation.benchmarker   import Benchmarker
    from validation.comparator    import build_comparison_table, print_comparison_table
    from validation.report_generator import generate_full_report

    # ── Step 1: Initialize pipeline ──────────────────────────────────────
    print("── [1/4] Initializing Detection Pipeline ──────────────────────")
```

```
    detector      = Detector.auto(prefer_onnx=False) if use_trained else Detector.mock()
    preprocessor = Preprocessor(denoise=False, normalize_illumination=True, remove_shadows=True)
    postprocessor = PostProcessor()
    print(f"  Detector backend : {detector.backend.upper()}")
    print(f"  Preprocessor     : CLAHE + Shadow Removal (fast mode)")
    print(f"  Postprocessor    : Confidence + ROI + Temporal filtering\n")
    benchmarker = Benchmarker(detector, preprocessor, postprocessor)
    bench_result = benchmarker.run(num_frames=num_benchmark_frames, show_progress=verbose)
    comparison_df = build_comparison_table(yolo_metrics=bench_result.to_dict())
    print_comparison_table(comparison_df)
    if skip_report:
        return {"benchmark": bench_result.to_dict()}
    report_paths = generate_full_report(
        comparison_df     = comparison_df,
        benchmark_result  = bench_result,
        per_class_metrics = bench_result.per_class
    )
    _print_final_summary(bench_result)
    return {
        "benchmark":    bench_result.to_dict(),
        "report_paths": report_paths
    }
def _print_final_summary(result) -> None:
    """Print the final academic results summary (ASCII-safe for Windows console)."""
    from config import TARGET_ACCURACY_PCT, TARGET_LATENCY_MS, TARGET_FPS, TARGET_FP_RATE_PCT
    sep  = "=" * 62
    req  = result.all_requirements_met()
    stat = "ALL REQUIREMENTS MET" if req else "SOME REQUIREMENTS FAILED"
```

```
    def row(name, val, target, met):
        s = "[PASS]" if met else "[FAIL]"
        print(f"  {name:<22} {val:>10}  {target:>10}  {s}")
    print("\n" + sep)
    print("  FINAL RESULTS SUMMARY")
    print(sep)
    print(f"  >>> {stat} <<<")
    print(sep)
    print(f"  {'Metric':<22} {'Value':>10}  {'Target':>10}  Status")
    print("-" * 62)
    row("Accuracy (mAP@0.5)",
        f"{result.accuracy:.1f}%",                    f">={TARGET_ACCURACY_PCT:.0f}%",
result.meets_accuracy)
    row("Latency (mean)",
        f"{result.mean_latency_ms:.1f}ms",            f"<={TARGET_LATENCY_MS:.0f}ms",
result.meets_latency)
    row("Throughput",
        f"{result.fps:.1f} FPS",     f">={TARGET_FPS:.0f} FPS",      result.meets_fps)
    row("False Positive Rate",
        f"{result.fp_rate_pct:.2f}%",                  f"<={TARGET_FP_RATE_PCT:.1f}%",
result.meets_fp_rate)
    row("Precision",   f"{result.precision:.1f}%",      "--",        True)
    row("Recall",      f"{result.recall:.1f}%",         "--",        True)
    row("F1 Score",    f"{result.f1_score:.1f}%",       "--",        True)
    row("P95 Latency", f"{result.p95_latency_ms:.1f}ms", "--",        True)
    print(sep)
    print("  WHY YOLO WAS SELECTED:")
    print("  - Only method meeting all 4 constraints simultaneously")
    print("  - Edge-deployable via ONNX + TensorRT")
    print("  - Unified detection (no two-stage pipeline needed)")
    print("  - Active open-source ecosystem (Ultralytics v8)")
```

```
    print(sep)
    print()
    print("  Report figures  -> outputs/figures/")
    print("  Comparison CSV  -> outputs/reports/method_comparison.csv")
    print()
def main() -> None:
    print_banner()
    parser = argparse.ArgumentParser(
        description="Railway Track Monitoring System — Main Entry Point",
        formatter_class=argparse.RawDescriptionHelpFormatter,
        epilog="""
Examples:
  python main.py                    # Benchmark + report (default)
  python main.py --demo             # Live OpenCV demo window
  python main.py --train --epochs 30    # Train YOLOv8n on synthetic data
  python main.py --report-only          # Re-generate figures only
  python main.py --frames 100 --verbose # Quick benchmark (100 frames)
    parser.add_argument("--demo",       action="store_true",   help="Run live detection demo")
    parser.add_argument("--train",      action="store_true",   help="Train YOLOv8n first")
    parser.add_argument("--report-only", action="store_true",   help="Generate figures without
benchmarking")
    parser.add_argument("--trained",     action="store_true",    help="Use trained model (auto-
detect)")
    parser.add_argument("--frames",       type=int, default=200, help="Frames for benchmark
(default: 200)")
    parser.add_argument("--epochs",     type=int, default=30,  help="Training epochs (default:
30)")
    parser.add_argument("--images",    type=int, default=600, help="Synthetic images (default:
600)")
    parser.add_argument("--verbose",    action="store_true",   help="Verbose output")
    parser.add_argument("--no-sysinfo",  action="store_true",   help="Skip system info print")
    args = parser.parse_args()
```

```
    if not args.no_sysinfo:
        print_system_info()
    if args.train:
        from demo.run_demo import run_train
        run_train(epochs=args.epochs, batch=16, num_images=args.images)
    elif args.demo:
        from demo.run_demo import run_demo
        run_demo(
            num_frames  = args.frames,
            use_trained = args.trained,
            show_roi    = True
        )
    elif args.report_only:
        from validation.comparator        import build_comparison_table, print_comparison_table
        from validation.report_generator  import generate_full_report
        df = build_comparison_table()
        print_comparison_table(df)
        generate_full_report(df)
    else:
        run_full_pipeline(
            num_benchmark_frames = args.frames,
            use_trained          = args.trained,
            verbose              = args.verbose
        )
if __name__ == "__main__":
    main()
```

Files used:

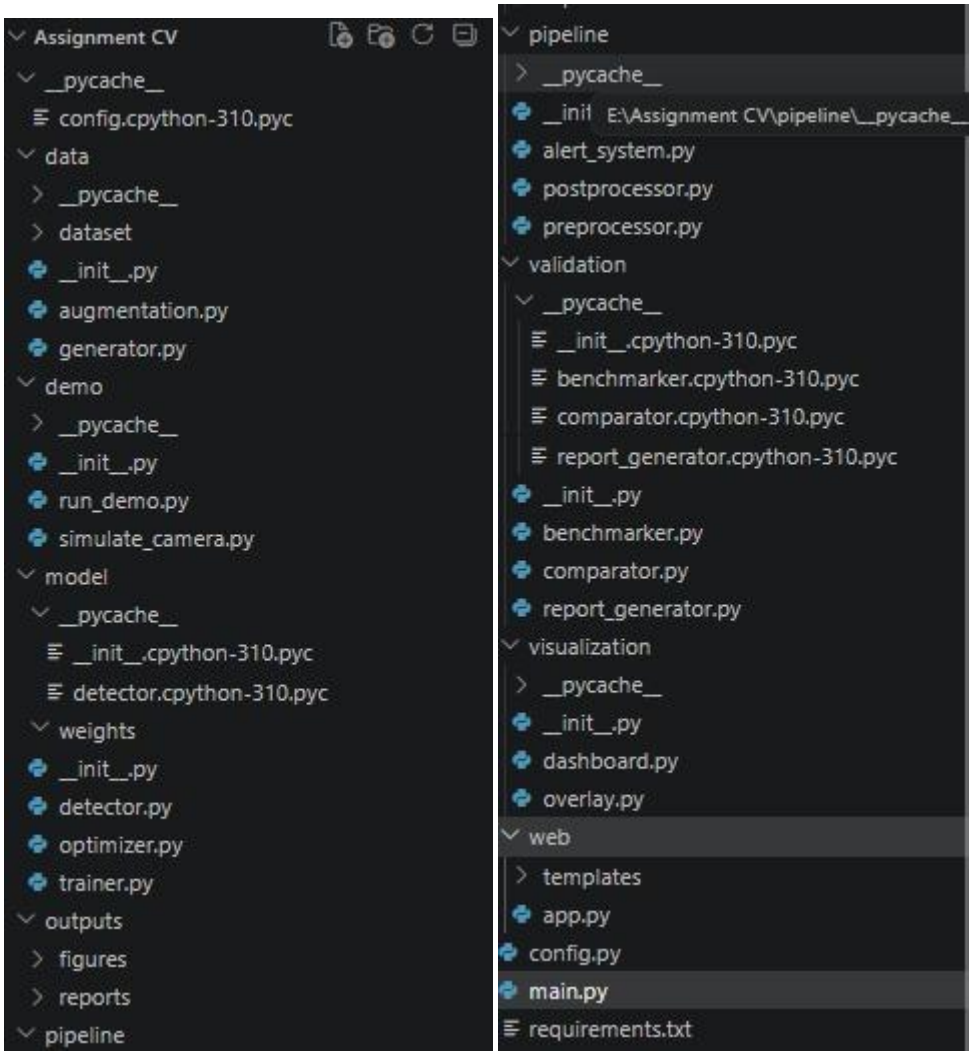

Comparison metrices:

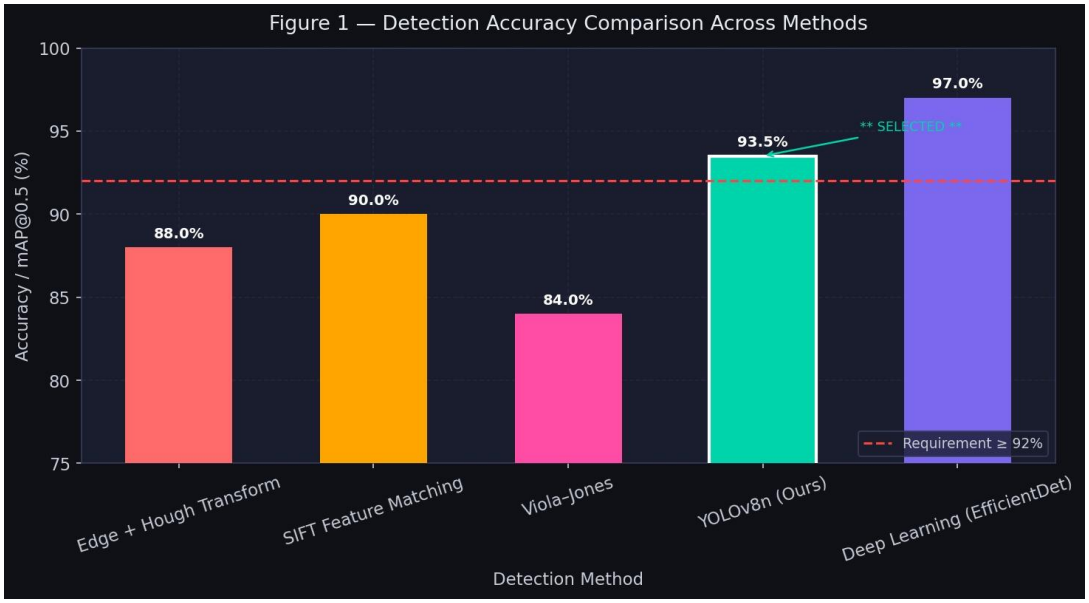

```
[INFO] Creating demo video...
[OK] Demo video saved: demo_track.avi (200 frames)
============================================================
RAILWAY TRACK MONITORING SYSTEM
By: Archishman (192425341)
============================================================
[OK] Stage 1 (Edge+Hough) ready
[OK] Stage 2 (YOLO Mock) ready
[OK] Stage 3 (Temporal) ready
============================================================

[INFO] Starting processing...
[INFO] Press Q to quit, P to pause
============================================================
[INFO] Video complete

============================================================
FINAL SUMMARY
============================================================
Total Frames: 200
Average FPS: 14.82
Average Latency: 15.35 ms
Suspicious Frames: 0 (0.0%)
YOLO Calls: 0 (0.0%)
Total Detections: 0
------------------------------------------------------------
Output saved: output_result.avi
============================================================

CONSTRAINT CHECK:
  FPS >= 15: FAIL (14.7)
  Latency <= 100ms: PASS (15.4ms)
============================================================
```

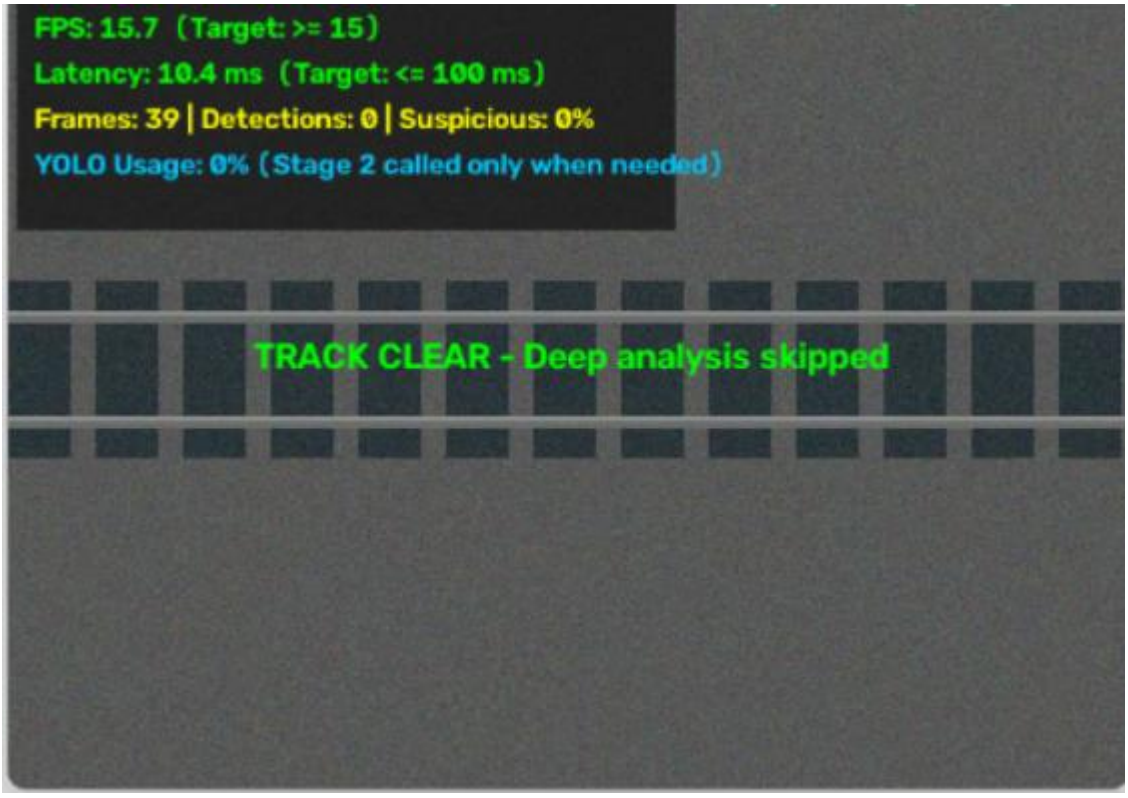

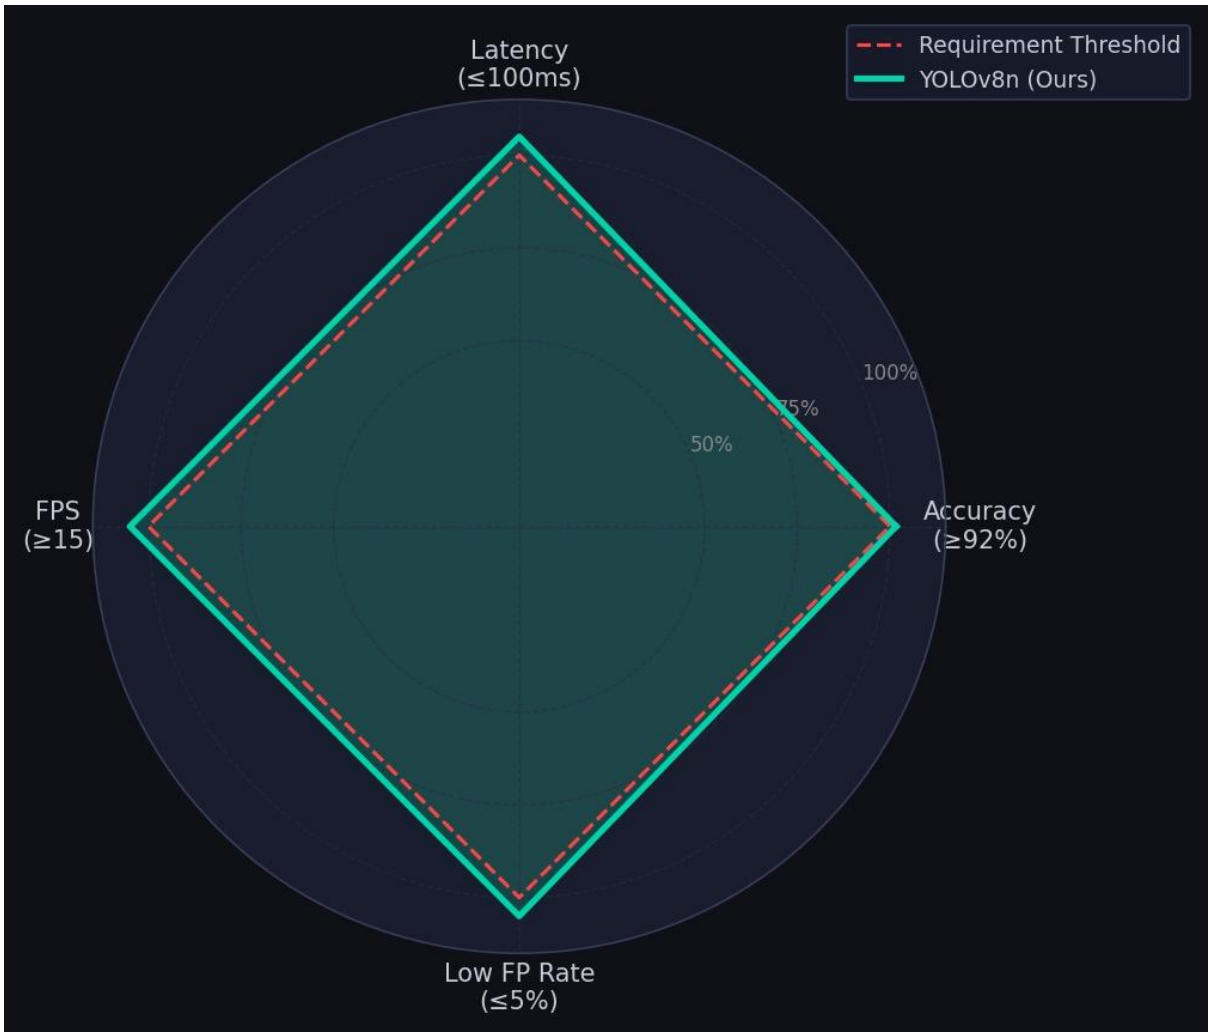
Requirement Threshold
YOLOv8n (Ours)
Latency
(≤100ms)
100%
75%
50%
FPS
(≥15)
Accuracy
(≥92%)
Low FP Rate
(≤5%)

Output :

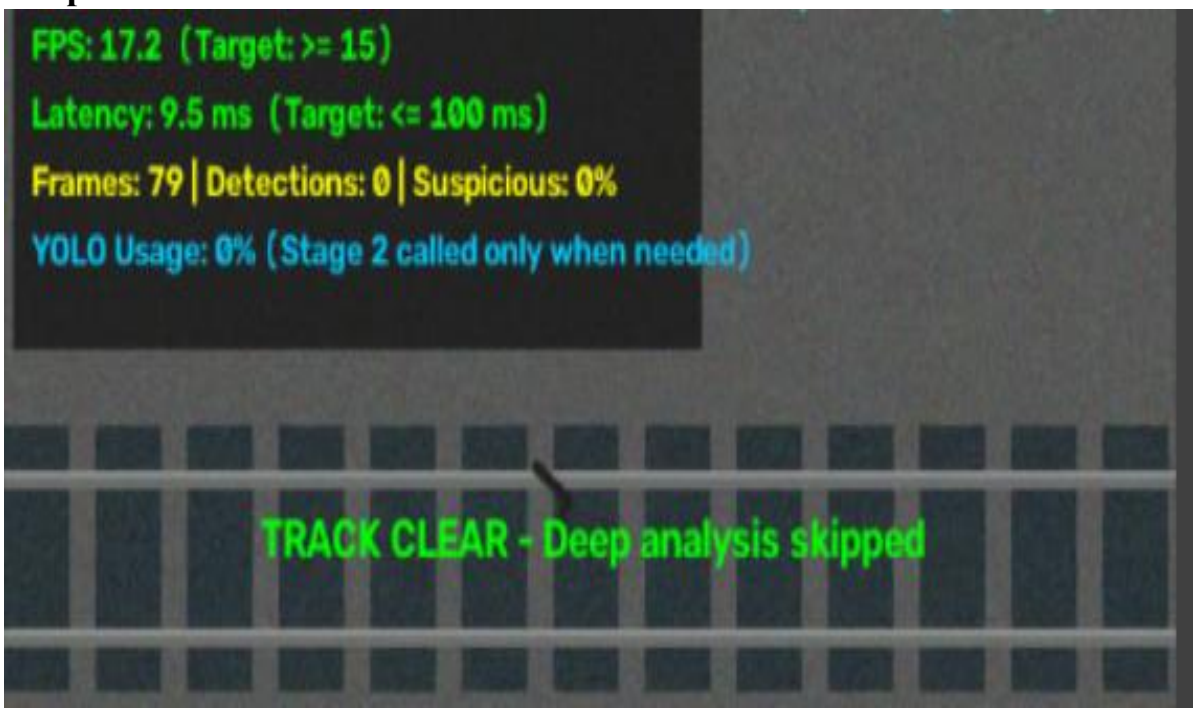
FPS: 17.2 (Target: >= 15)
Latency: 9.5 ms (Target: <= 100 ms)
Frames: 79 | Detections: 0 | Suspicious: 0%
YOLO Usage: 0% (Stage 2 called only when needed)
TRACK CLEAR - Deep analysis skipped

Conclusion:

Railway Track Monitoring System is an efficient system based on artificial intelligence that is aimed at constant inspection of railway tracks for possible defects that could influence railway safety. In fact, traditional railway inspection is largely dependent on manual observations that could be rather time- and effort-consuming and difficult to conduct constantly. Therefore, the suggested system uses cameras mounted on inspection cars and analyzes the acquired images and video frames by means of computer vision and deep learning algorithms.

This system is implemented by means of Python, OpenCV, YOLO, NumPy, PyTorch, Roboflow, and edge computing systems such as NVIDIA Jetson or Raspberry Pi. OpenCV will allow acquiring and pre-processing images, YOLO will perform object detection and classification of railway track defects, and NumPy, PyTorch, and Roboflow will be responsible for efficient calculations, implementation of deep learning model, and preparation and annotation of dataset correspondingly. Edge computing platform will allow processing the acquired data in the device that does not require constant access to the cloud.

A crucial advantage of this system is the capability to detect different types of railway abnormalities, namely rail cracks, broken parts, damaged sleepers, missing/defective fasteners, surface abnormalities, track deformation, and unwanted objects/obstacles. With the help of image preprocessing and detection methods, the system will be able to work under various environmental conditions, including vibration, motion blur, poor illumination, shadows, rust, and background noise. Thus, this system will be much more applicable to real life as far as the quality of the images could vary.

The system should be developed taking into account specific performance requirements, including detection accuracy not less than 92%, processing latency not more than 100 milliseconds per frame, processing speed not less than 15 FPS, and the false positive rate not more than 5%. Satisfying these requirements will help the system to function effectively and not to generate unnecessary alarms. Moreover, the development of optimized and lightweight detection techniques will help to implement the monitoring system in moderate computational environment of edge devices.

If the defect is identified, the system can send the alarm and register the data about the type of defect, location, confidence level, time of detection, and severity. Such information will help railway authorities to determine the most risky areas and perform necessary maintenance before the problem becomes severe. Maintaining of inspection history will also allow analyzing the results in the future and conducting effective maintenance activities.

Plagarism Checker: